

Applied AI

Lecture 2

Dr Artie Basukoski

[slides adapted from Artificial Intelligence: A Modern Approach, Russel and Norvig]

Agenda


- Agents and environments
- Types of agents
- Representing problems
- Selecting a state space
- Tree search algorithms
- Breadth first search
- Depth first search

- Next steps

Goal

- Identify the concept of an intelligent agent.
- Develop a small set of design principles for building successful agents.
- Agents should be rational – one that does the right thing.
- Behaviour depends on the environment and the goals that we define for the agents.
- We define a number of basic agent designs.

Agents and environments




Agents include humans, robots, machines, etc.

The **agent function** maps from percept histories to actions:

$$f : P^* \rightarrow A$$

The **agent program** runs on the physical

Vacuum-cleaner world



Percepts: location and contents, e.g., [A, *Dirty*]

Actions: *Left*, *Right*, *Suck*, *NoOp*

A vacuum-cleaner agent

Percept sequence	Action
<i>[A, Clean]</i>	<i>Right</i>
<i>[A, Dirty]</i>	<i>Suck</i>
<i>[B, Clean]</i>	<i>Left</i>
<i>[B, Dirty]</i>	<i>Suck</i>
<i>[A, Clean], [A, Clean]</i>	<i>Right</i>
<i>[A, Clean], [A, Dirty]</i>	<i>Suck</i>
<i>.</i>	<i>.</i>

```
function Reflex-Vacuum-Agent( [location,status]) returns an action
  if status = Dirty then return Suck
  else if location = A then return
    Right else if location = B then
    return Left
```

What is the **right** function?

Can it be implemented in a small agent

Rationality

Fixed **performance measure** evaluates the **environment sequence**

- one point per square cleaned up in time T ?
- one point per clean square per time step, minus one per move?
- penalize for $> k$ dirty squares?

A **rational agent** chooses whichever action maximizes the **expected** value of the performance measure **given the percept sequence to date**

Rational \neq omniscient

–percepts may not supply all relevant information Rational \neq clairvoyant

–action outcomes may not be as expected Hence, rational \neq successful

Rational \Rightarrow exploration, learning

PEAS

To design a rational agent, we must specify the **task environment**

Consider, e.g., the task of designing an automated taxi:

Performance measure: safety, destination, profits, legality, comfort, . . .

Environment: streets/freeways, traffic, pedestrians, weather, . . .

Actuators: steering, accelerator, brake, horn, speaker/display, . . .

Sensors: video, accelerometers, engine sensors, keyboard, GPS, . . .

Internet shopping agent

Performance measure: price, quality, appropriateness, efficiency

Environment: current and future WWW sites, vendors, shippers
Actuators: display to user, follow URL, fill in form

Sensors: HTML pages (text, graphics, scripts)

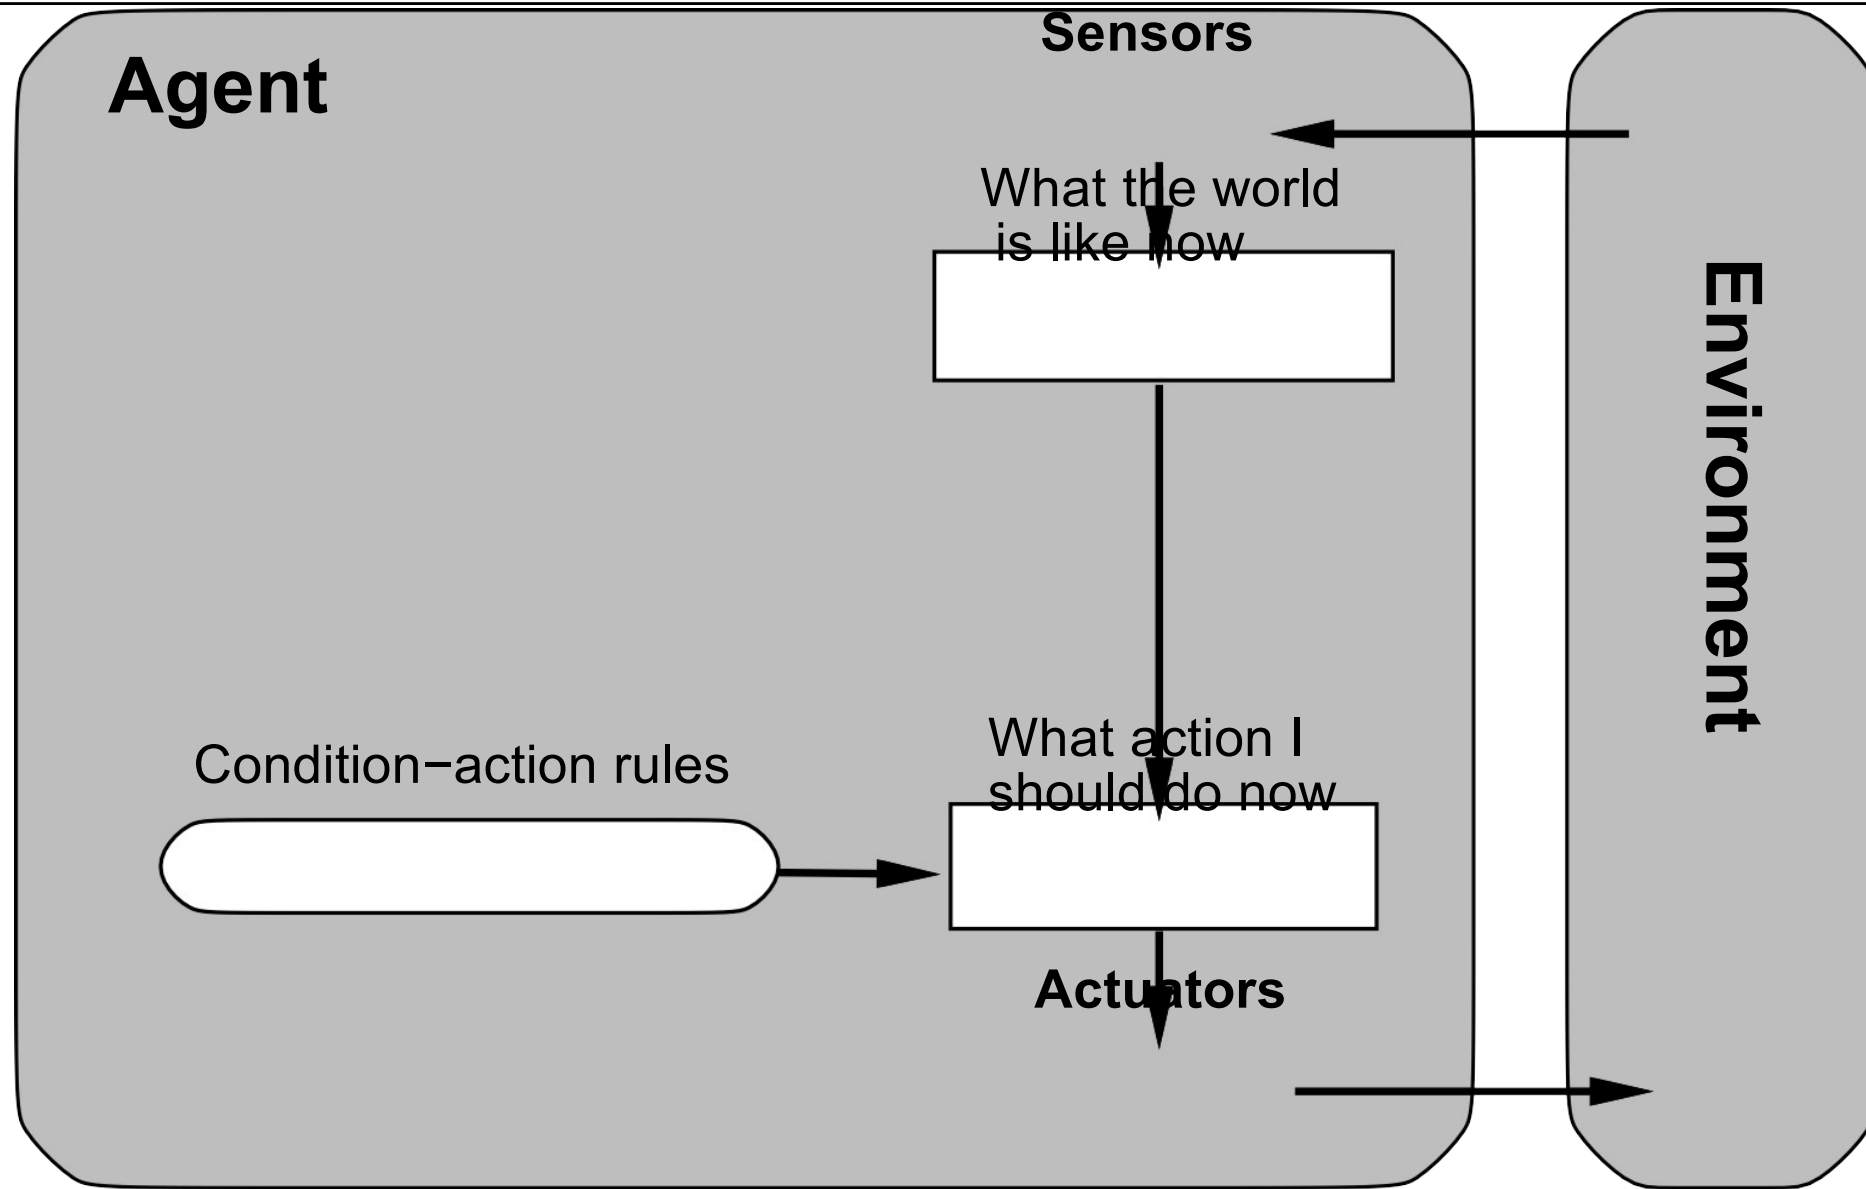
Environment types

Task Environment	Observable	Agents	Deterministic	Episodic	Static	Discrete
Crossword puzzle	Fully	Single	Deterministic	Sequential	Static	Discrete
Chess with a clock	Fully	Multi	Deterministic	Sequential	Semi	Discrete
Poker	Partially	Multi	Stochastic	Sequential	Static	Discrete
Backgammon	Fully	Multi	Stochastic	Sequential	Static	Discrete
Taxi driving	Partially	Multi	Stochastic	Sequential	Dynamic	Continuous
Medical diagnosis	Partially	Single	Stochastic	Sequential	Dynamic	Continuous
Image analysis	Fully	Single	Deterministic	Episodic	Semi	Continuous
Part-picking robot	Partially	Single	Stochastic	Episodic	Dynamic	Continuous
Refinery controller	Partially	Single	Stochastic	Sequential	Dynamic	Continuous
English tutor	Partially	Multi	Stochastic	Sequential	Dynamic	Discrete

The environment type largely determines the agent design

The real world is (of course) partially observable,
stochastic, sequential,
dynamic, continuous, multi-agent

Simple reflex agents



Example

```
function Reflex-Vacuum-Agent( [location,status]) returns an action
  if status = Dirty then return Suck
  else if location = A then return
    Right else if location = B then
    return Left
```

Use example from Jupyter lab

Create example from the vacuum environment

Problem-solving agents

Restricted form of general

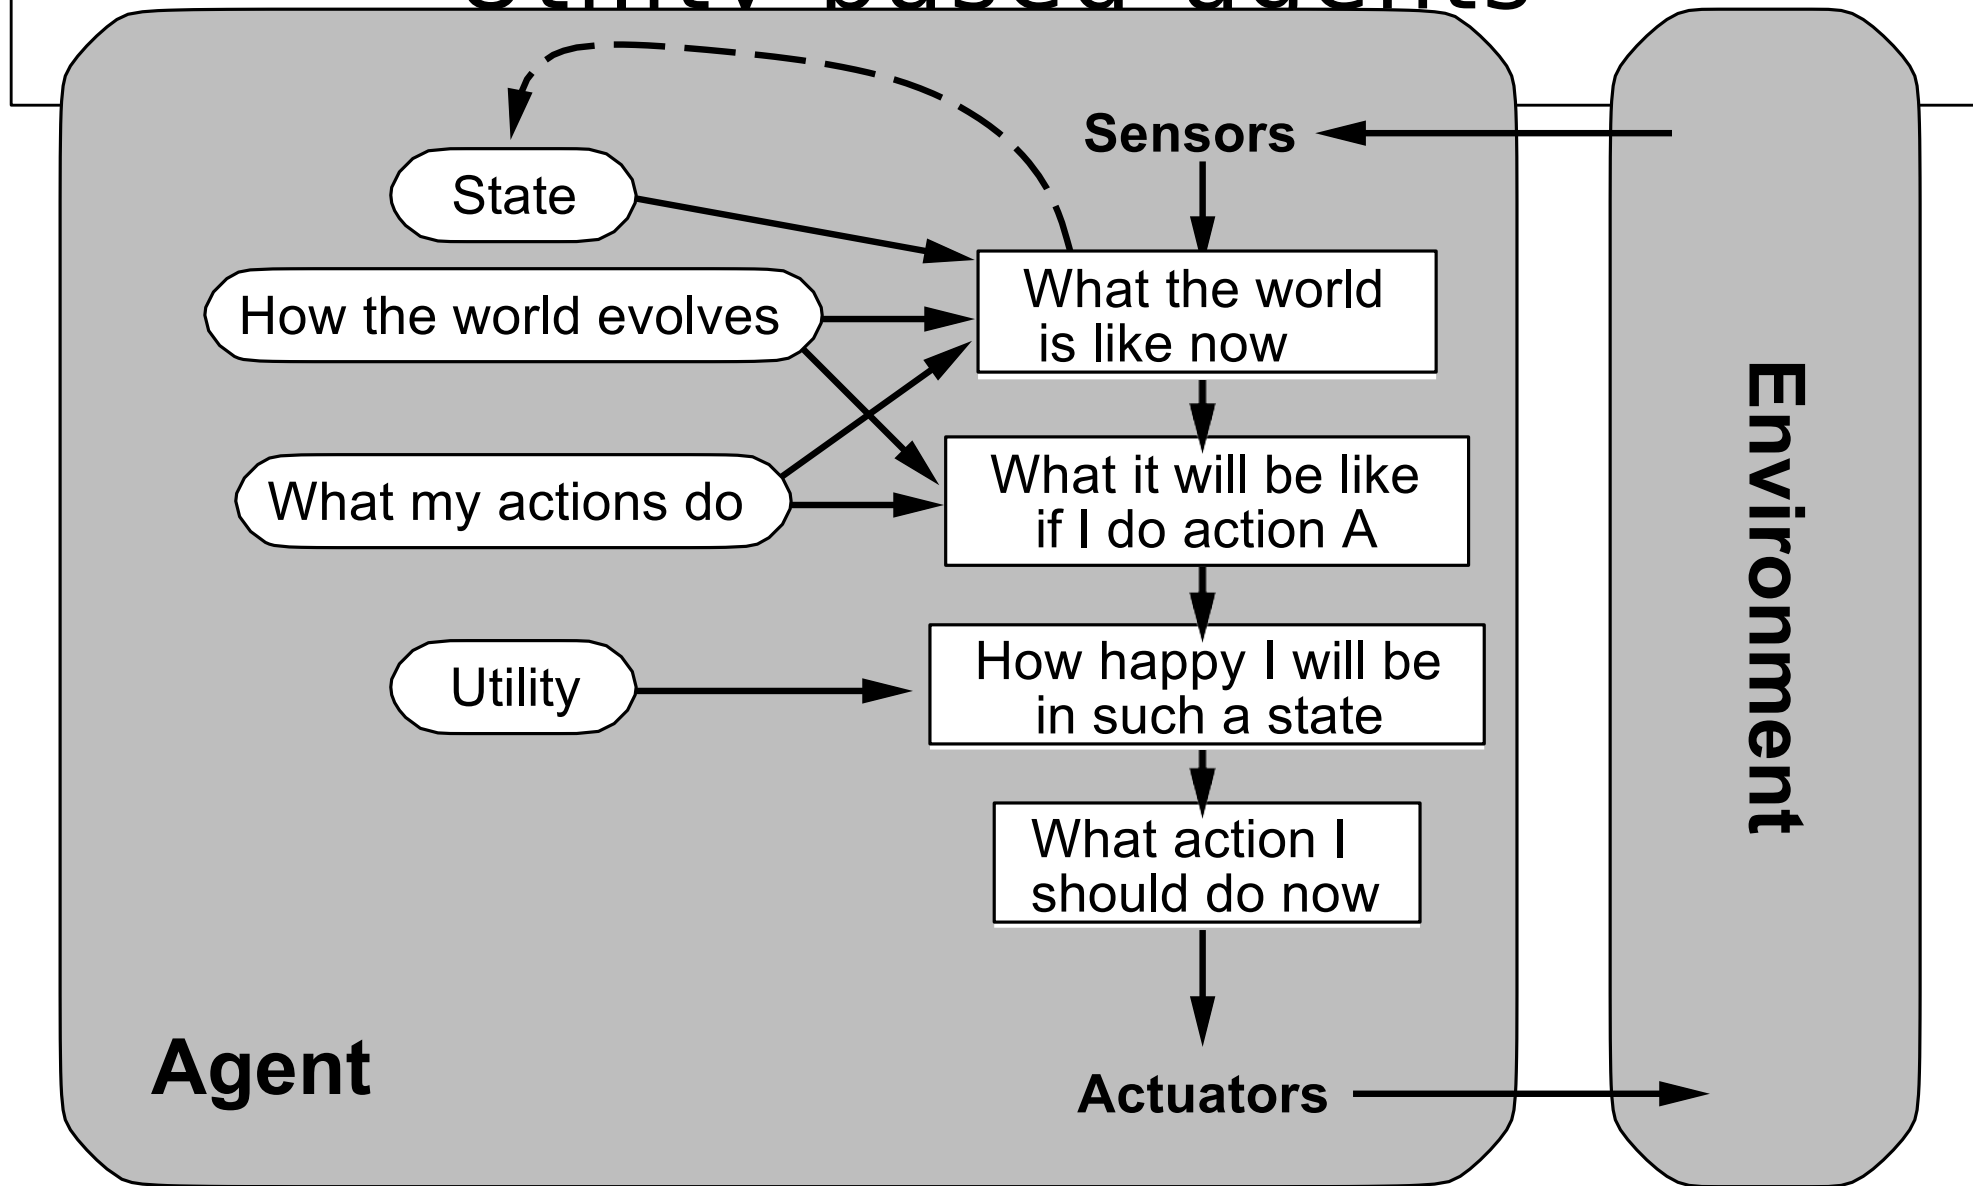
agent:

```
function Simple-Problem-Solving-Agent( percept) returns an action
  static: seq, an action sequence, initially empty
          state, some description of the current world state
          goal, a goal, initially null
          problem, a problem formulation

  state ← Update-State(state, percept)
  if seq is empty then
    goal ← Formulate-Goal(state)
    problem ← Formulate-Problem(state,
    goal) seq ← Search( problem)
  action ← Recommendation(seq,
  state) seq ← Remainder(seq,
  state)
  return action
```

Note: this is **offline** problem solving; solution executed

Utility-based agents



Summary

Agents interact with environments through actuators and sensors

The agent function describes what the agent does in all circumstances The performance measure evaluates the environment sequence

A perfectly rational agent maximizes expected performance Agent programs implement (some) agent functions

PEAS descriptions define task environments

Environments are categorized along several dimensions:

Example: Romania

On holiday in Romania; currently in Arad. Flight leaves tomorrow from Bucharest

Formulate goal:

be in Bucharest

Formulate problem:

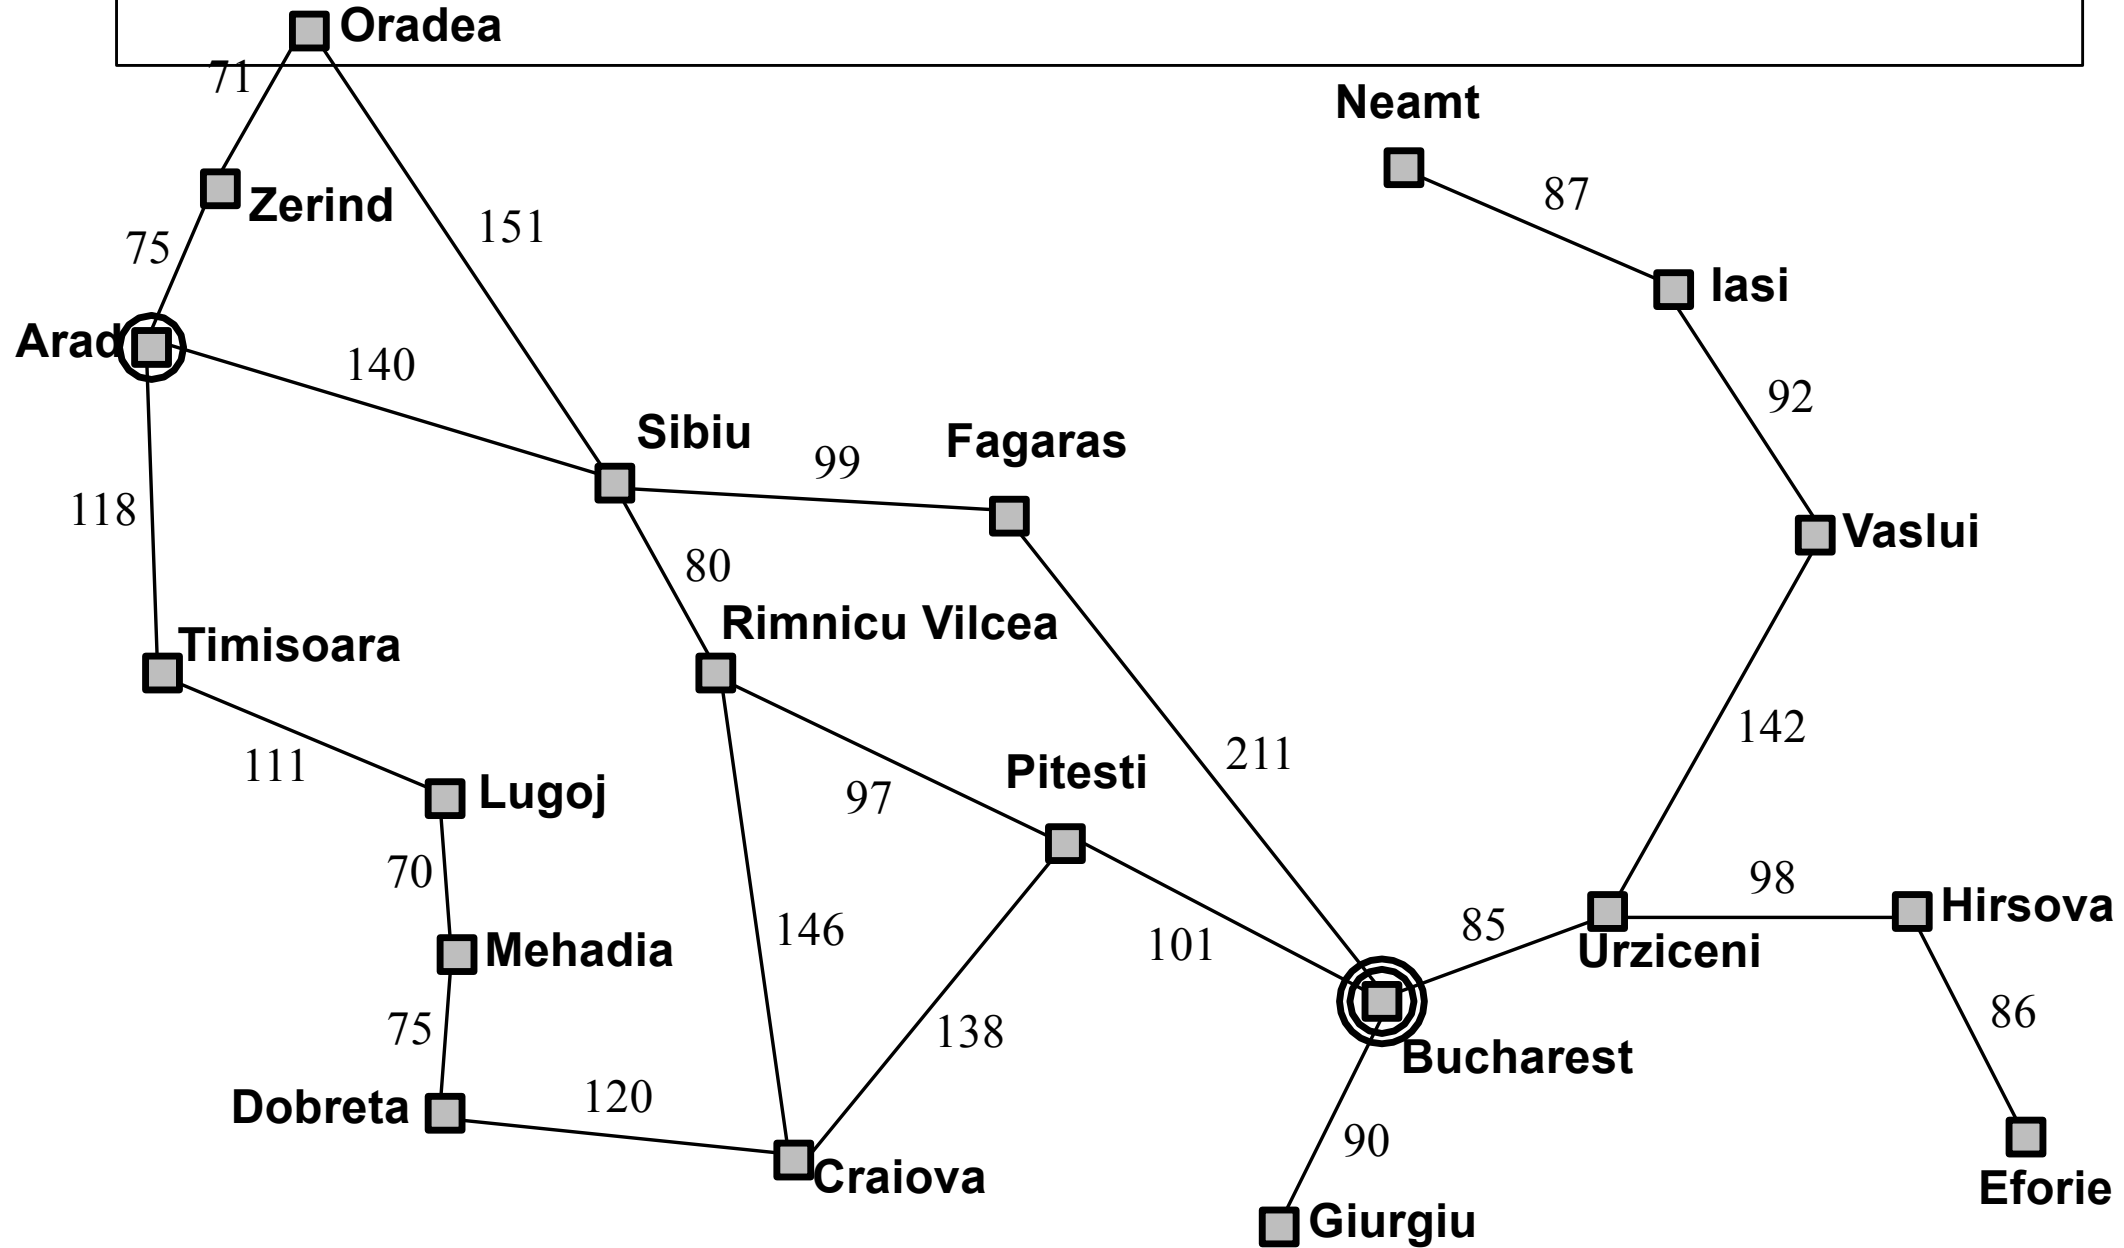
states: various cities

actions: drive between cities

Find solution:

sequence of cities, e.g.,
Arad, Sibiu, Fagaras,
Bucharest

Example: Romania



Problem types

~~Deterministic, fully observable~~ \Rightarrow ~~single-state problem~~

Agent knows exactly which state it will be in; solution is a sequence

Non-observable \Rightarrow conformant problem

Agent may have no idea where it is; solution (if any) is a sequence

Nondeterministic and/or partially observable \Rightarrow
contingency problem percepts provide new
information about current state

solution is a contingent plan or a policy
often **interleave** search,
execution

Unknown state space \Rightarrow exploration
problem (“online”)

Single-state problem formulation

A **problem** is defined by four items:

initial state e.g., “at Arad”

successor function $S(x)$ = set of action-state pairs

e.g., $S(\text{Arad}) = \{(\text{Arad} \rightarrow \text{Zerind}, \text{Zerind}), \dots\}$

goal test, can be

explicit, e.g., $x = \text{“at Bucharest”}$ **implicit**, e.g., $\text{NotDirt}(x)$

path cost (additive)

e.g., sum of distances,
number of actions executed,
etc.

$c(x, a, y)$ is the **step cost**

Selecting a state space

Real world is absurdly complex

⇒ state space must be **abstracted** for problem solving


(Abstract) state = set of real states

(Abstract) action = complex combination of real actions e.g., “Arad → Zerind” represents a complex set

of possible routes, detours, rest stops, etc. For guaranteed realizability, **any** real state “in Arad” must get to **some** real state “in Zerind”


(Abstract) solution =
set of real paths that are

Example: vacuum world state space graph



states: integer dirt and robot locations (ignore dirt amounts etc.) **actions:** *Left*, *Right*, *Suck*, *NoOp*
goal test: no dirt
path cost: 1 per action (0 for *NoOp*)


Example: The 8-puzzle



Start State
Goal State

states: integer locations of tiles (ignore intermediate positions) **actions**: move blank left, right, up, down (ignore unjamming etc.) **goal test** = goal state (given)
path cost: 1 per move

Example: robotic assembly



states: real-valued coordinates of robot joint angles
parts of the object to be assembled

actions: continuous motions of robot joints

goal test: complete assembly **with no robot included!**

path cost: time to execute

Tree search algorithms

Basic idea:

offline, simulated exploration of state
space

by generating successors of already-
explored states (a.k.a. **expanding**

```
function TreeSearch( problem, strategy ) returns a solution,  
or failure  
    initialize the search tree using the initial state of  
    problem
```

```
    loop do
```

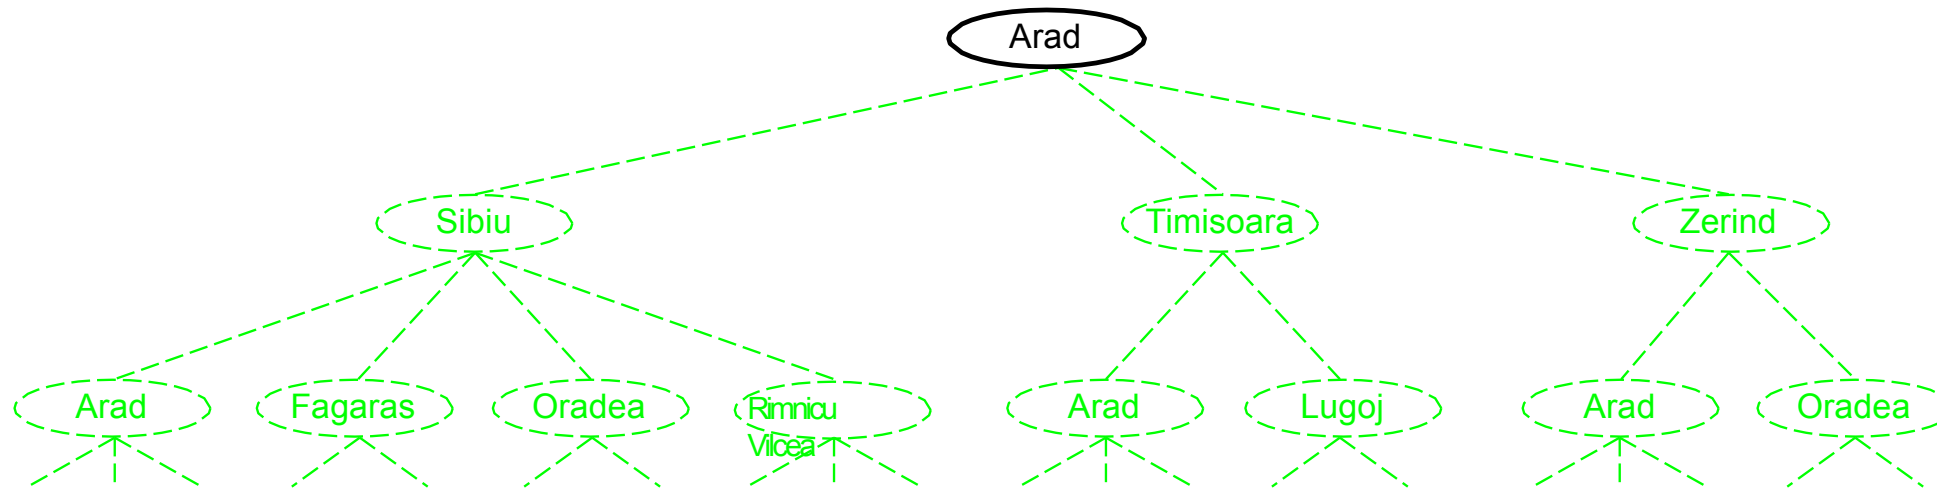
```
        if there are no candidates for expansion then return  
            failure  
        choose a leaf node for expansion according to  
        strategy
```

```
        if the node contains a goal state then return the  
            corresponding solution
```

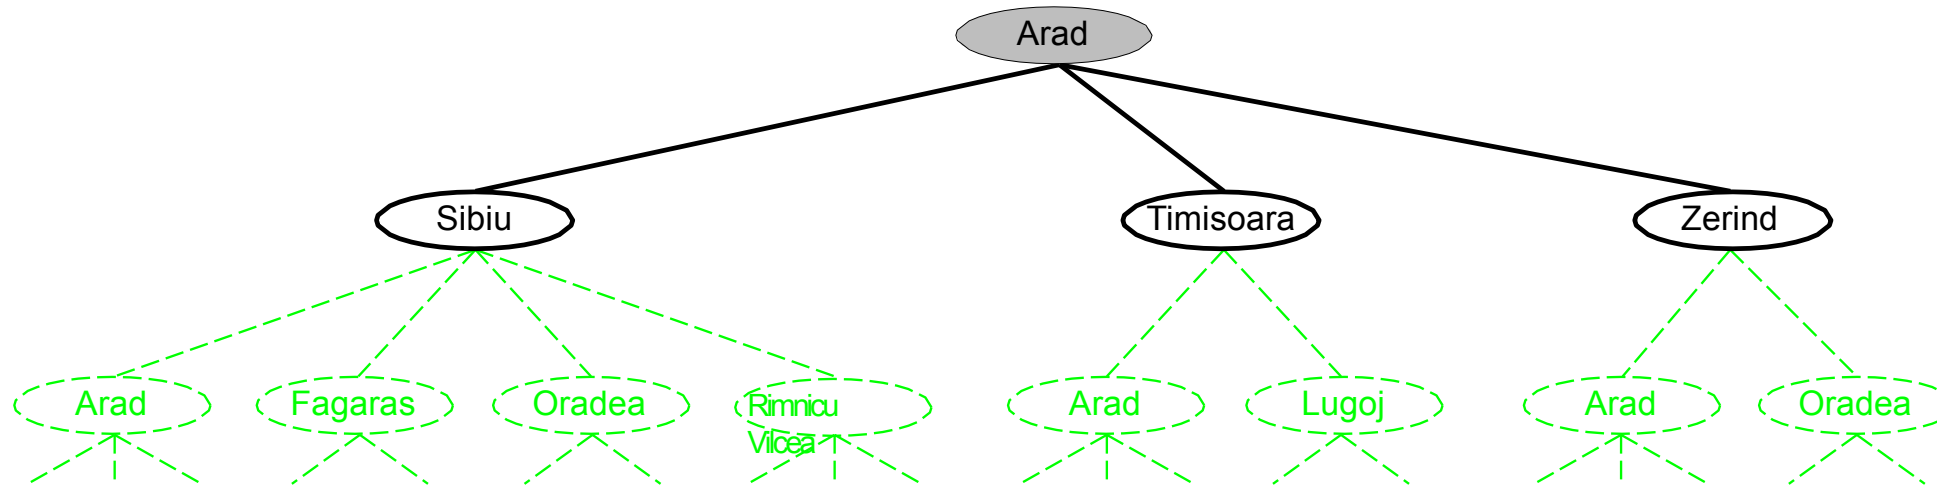
```
        else expand the node and add the resulting nodes to  
            the search tree
```

```
    end
```

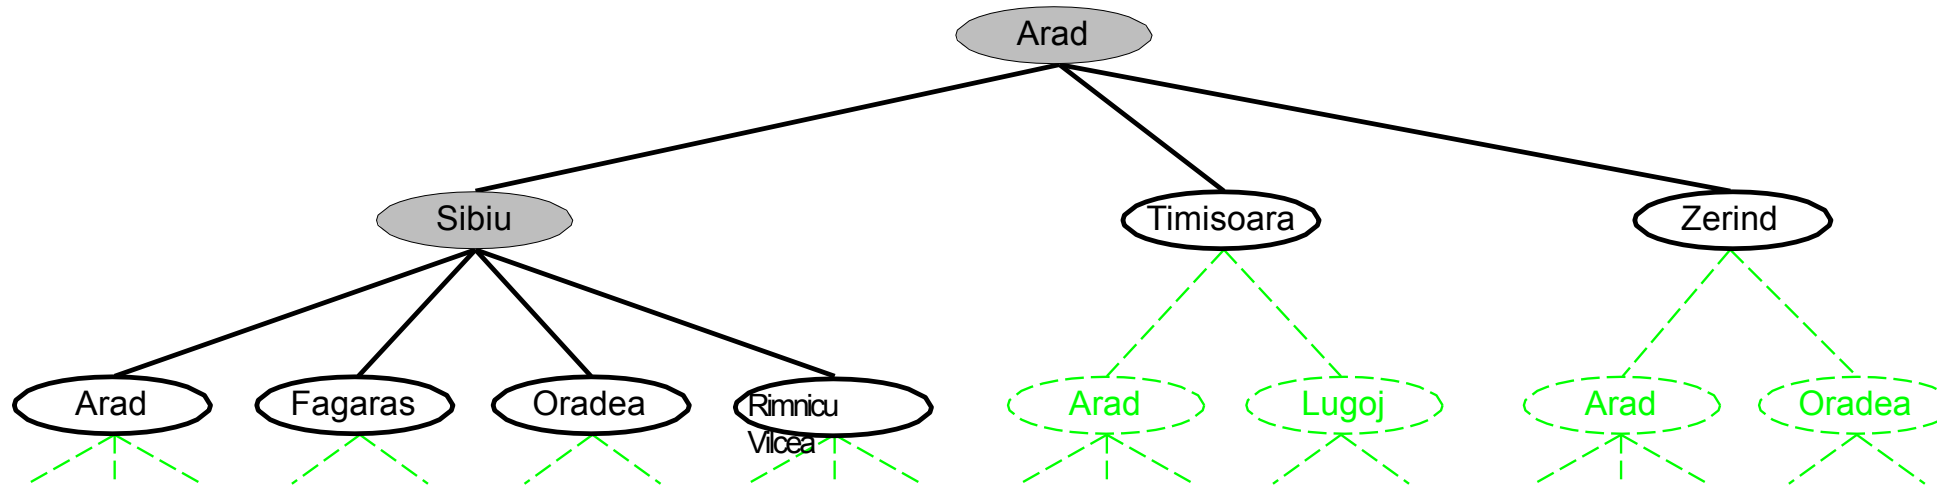

Tree search example



Tree search example



Tree search example




Implementation: states vs. nodes

A **state** is a (representation of) a physical configuration

A **node** is a data structure constituting part of a search tree includes **parent**, **children**, **depth**, **path cost** $g(x)$

States do not have parents, children, depth, or path cost!



The `Expand` function creates new nodes, filling in the various fields and using the `SuccessorFn` of the problem to create the corresponding states.

Implementation: general tree search

```
function Tree-Search( problem, fringe) returns a solution,  
or failure fringe ← Insert(Make-Node(Initial-  
State[problem]), fringe) loop do  
  if fringe is empty then return failure  
  node ← Remove-Front(fringe)  
  if Goal-Test(problem, State(node)) then return  
  node fringe ← InsertAll(Expand(node,  
  problem), fringe)
```

```
function Expand( node, problem) returns a set of nodes  
successors ← the empty set  
for each action, result  
in Successor-  
Fn(problem,  
State[node]) do  
  s ← a new Node  
  Parent-Node[s] ←  
  node;
```

Search strategies

A strategy is defined by picking the **order of node expansion**

Strategies are evaluated along the following dimensions:

completeness—does it always find a solution if one exists? **time complexity**—number of nodes generated/expanded **space complexity**—maximum number of nodes in memory **optimality**—does it always find a least-cost solution?

Time and space complexity are measured in terms of **b** —maximum branching factor of the search tree **d** —depth of the least-cost solution

m —maximum depth of the state space (may be ∞)

Uninformed search strategies

Uninformed strategies use only the information available in the problem definition


Breadth-first search Uniform-cost
search Depth-first search Depth-
limited search
Iterative deepening search

Breadth-first search

Expand shallowest unexpanded node

Implementation:

fringe is a FIFO queue, i.e., new successors go at end




Breadth-First Search

Expand shallowest unexpanded node

Implementation:

fringe is a FIFO queue, i.e., new successors go at end




Breadth-First Search

Expand shallowest unexpanded node

Implementation:

fringe is a FIFO queue, i.e., new successors go at end




Breadth-First Search

Expand shallowest unexpanded node

Implementation:

fringe is a FIFO queue, i.e., new successors go at end




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front



Questions Discussion?